

# Kubernetes Cluster Setup Guide - Create Your Own K8s Cluster

## Overview

This guide will help you create a production-ready Kubernetes cluster using **kubeadm** with **containerd** as the container runtime. We'll set up a multi-node cluster with proper networking and security.

## Prerequisites

### System Requirements

- **Minimum 2 GB RAM** per node (4 GB recommended)
- **2 CPUs** per node
- **Network connectivity** between all nodes
- **Unique hostname, MAC address, and product\_uuid** for every node

### Supported Operating Systems

- Ubuntu 16.04+ (Recommended: Ubuntu 20.04/22.04)
  - CentOS 7+
  - RHEL 7+
  - Debian 9+
- 

## Phase 1: Environment Setup (All Nodes)

### Step 1: Update System & Install Dependencies

```
bash

# Update package index
sudo apt update && sudo apt upgrade -y

# Install required packages
sudo apt install -y apt-transport-https ca-certificates curl software-properties-common

# Install additional tools
sudo apt install -y wget gpg lsb-release
```

### Step 2: Disable Swap (Critical for Kubernetes)

```
bash
```

```
# Disable swap immediately
```

```
sudo swapoff -a
```

```
# Disable swap permanently
```

```
sudo sed -i '/ swap / s/^(.*)$/# \1/g' /etc/fstab
```

```
# Verify swap is disabled
```

```
free -h
```

### Step 3: Configure Kernel Modules & Networking

```
bash
```

```
# Load required kernel modules
```

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
```

```
overlay
```

```
br_netfilter
```

```
EOF
```

```
sudo modprobe overlay
```

```
sudo modprobe br_netfilter
```

```
# Configure sysctl parameters
```

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOF
```

```
# Apply sysctl parameters
```

```
sudo sysctl --system
```

### Step 4: Configure Firewall (If Enabled)

```
bash
```

```
# For Ubuntu/Debian with ufw
```

```
sudo ufw allow 6443/tcp    # Kubernetes API server  
sudo ufw allow 2379:2380/tcp # etcd server client API  
sudo ufw allow 10250/tcp   # Kubelet API  
sudo ufw allow 10259/tcp   # kube-scheduler  
sudo ufw allow 10257/tcp   # kube-controller-manager  
sudo ufw allow 30000:32767/tcp # NodePort Services
```

```
# For worker nodes, also allow:
```

```
sudo ufw allow 10250/tcp   # Kubelet API  
sudo ufw allow 30000:32767/tcp # NodePort Services
```

---

## Phase 2: Install Container Runtime (containerd)

### Step 5: Install containerd

```
bash
```

```
# Add Docker's official GPG key
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyr
```

```
# Add Docker repository
```


```
echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyri
```

```
# Update package index
```

```
sudo apt update
```

```
# Install containerd
```

```
sudo apt install -y containerd.io
```



### Step 6: Configure containerd

```
bash
```

```
# Create containerd configuration directory
```

```
sudo mkdir -p /etc/containerd
```

```
# Generate default configuration
```

```
sudo containerd config default | sudo tee /etc/containerd/config.toml
```

```
# Enable SystemdCgroup (Important for kubeadm)
```

```
sudo sed -i 's/SystemdCgroup = false/SystemdCgroup = true/' /etc/containerd/config.toml
```

```
# Restart and enable containerd
```

```
sudo systemctl restart containerd
```

```
sudo systemctl enable containerd
```

```
# Verify containerd is running
```

```
sudo systemctl status containerd
```

---

## Phase 3: Install Kubernetes Components

### Step 7: Add Kubernetes Repository

```
bash
```

```
# Add Kubernetes GPG key
```

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc
```

```
# Add Kubernetes repository
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/st
```

```
# Update package index
```

```
sudo apt update
```



### Step 8: Install Kubernetes Tools

```
bash
```

```
# Install kubelet, kubeadm, and kubectl
```

```
sudo apt install -y kubelet kubeadm kubectl
```

```
# Hold packages to prevent automatic updates
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

```
# Enable kubelet service
```

```
sudo systemctl enable kubelet
```

---

## Phase 4: Initialize Master/Control Plane Node

### Step 9: Initialize Kubernetes Cluster (Master Node Only)

```
bash
```

```
# Initialize the cluster
```

```
sudo kubeadm init \  
  --pod-network-cidr=192.168.0.0/16 \  
  --kubernetes-version=stable \  
  --node-name=$(hostname -s)
```

```
# IMPORTANT: Save the join command output for worker nodes!
```

### Step 10: Configure kubectl (Master Node Only)

```
bash
```

```
# Set up kubectl for regular user
```

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
# Verify cluster status
```

```
kubectl get nodes
```

```
kubectl get pods -A
```

### Step 11: Install Pod Network Add-on (Master Node Only)

```
bash
```

```
# Install Calico CNI (Recommended)
```

```
kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v3.28.1/manifests/tigera
```

```
# Download and apply Calico custom resources
```

```
curl https://raw.githubusercontent.com/projectcalico/calico/v3.28.1/manifests/custom-resources.
```

```
# Apply Calico configuration
```

```
kubectl create -f custom-resources.yaml
```

```
# Wait for all pods to be ready
```

```
kubectl get pods -n calico-system --watch
```



## Alternative CNI Options:

```
bash
```

```
# Flannel (Simpler option)
```

```
kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yaml
```

```
# Weave Net
```

```
kubectl apply -f https://github.com/weaveworks/weave/releases/download/v2.8.1/weave-daemonset-k8s.yaml
```



## Phase 5: Join Worker Nodes

### Step 12: Join Worker Nodes to Cluster

```
bash
```

```
# On each worker node, run the join command from kubeadm init output:
```

```
sudo kubeadm join <MASTER-IP>:6443 \
  --token <TOKEN> \
  --discovery-token-ca-cert-hash sha256:<HASH>
```

```
# If you lost the join command, generate a new one on master:
```

```
kubeadm token create --print-join-command
```

### Step 13: Verify Cluster Setup

```
bash
```

```
# On master node, check all nodes
```

```
kubectl get nodes -o wide
```

```
# Check system pods
```

```
kubectl get pods -n kube-system
```

```
# Check cluster info
```

```
kubectl cluster-info
```

---

## Phase 6: Post-Installation Configuration

### Step 14: Remove Taint from Master (Optional - for single node)

```
bash
```

```
# Allow scheduling pods on master node (for development/testing)
```

```
kubectl taint nodes --all node-role.kubernetes.io/control-plane-
```

### Step 15: Install Kubernetes Dashboard (Optional)

```
bash
```

```
# Deploy Dashboard
```

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.7.0/aio/deploy/recon
```

```
# Create admin user
```

```
cat <<EOF | kubectl apply -f -
```

```
apiVersion: v1
```

```
kind: ServiceAccount
```

```
metadata:
```

```
  name: admin-user
```

```
  namespace: kubernetes-dashboard
```

```
---
```

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: ClusterRoleBinding
```

```
metadata:
```

```
  name: admin-user
```

```
roleRef:
```

```
  apiGroup: rbac.authorization.k8s.io
```

```
  kind: ClusterRole
```

```
  name: cluster-admin
```

```
subjects:
```

```
- kind: ServiceAccount
```

```
  name: admin-user
```

```
  namespace: kubernetes-dashboard
```

```
EOF
```

```
# Get access token
```

```
kubectl -n kubernetes-dashboard create token admin-user
```

```
# Access dashboard
```

```
kubectl proxy
```

```
# Visit: http://localhost:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes
```



## Step 16: Install Helm (Package Manager)



```
bash
```

```
# Install Helm
```

```
curl https://baltocdn.com/helm/signing.asc | gpg --dearmor | sudo tee /usr/share/keyrings/helm.  
echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/helm.gpg] https://ba  
sudo apt update  
sudo apt install helm
```

```
# Verify Helm installation
```

```
helm version
```

---

## Testing Your Cluster

### Step 17: Deploy Test Application

```
bash
```

```
# Create a test deployment
```

```
kubectl create deployment nginx --image=nginx:latest
```

```
# Expose the deployment
```

```
kubectl expose deployment nginx --port=80 --type=NodePort
```

```
# Get service details
```

```
kubectl get services
```

```
# Test the application
```

```
curl http://<NODE-IP>:<NODE-PORT>
```

---

## Cluster Management Commands

### Essential kubectl Commands

```
bash
```

```
# Node management
```

```
kubectl get nodes
```

```
kubectl describe node <NODE-NAME>
```

```
kubectl drain <NODE-NAME> --ignore-daemonsets
```

```
kubectl uncordon <NODE-NAME>
```

```
# Pod management
```

```
kubectl get pods -A
```

```
kubectl logs <POD-NAME> -n <NAMESPACE>
```

```
kubectl exec -it <POD-NAME> -- /bin/bash
```

```
# Cluster health
```

```
kubectl get componentstatuses
```

```
kubectl top nodes
```

```
kubectl top pods -A
```

## Cluster Maintenance

```
bash
```

```
# Backup etcd (Important!)
```

```
ETCDCTL_API=3 etcdctl snapshot save backup.db \  
  --endpoints=https://127.0.0.1:2379 \  
  --cacert=/etc/kubernetes/pki/etcd/ca.crt \  
  --cert=/etc/kubernetes/pki/etcd/server.crt \  
  --key=/etc/kubernetes/pki/etcd/server.key
```

```
# Upgrade cluster
```

```
sudo kubeadm upgrade plan
```

```
sudo kubeadm upgrade apply v1.31.x
```

## Troubleshooting

### Common Issues & Solutions

#### 1. Node Not Ready

```
bash
```

```
# Check kubelet Logs
```

```
sudo journalctl -u kubelet -f
```

```
# Check node conditions
```

```
kubectl describe node <NODE-NAME>
```

## 2. Pods Stuck in Pending

```
bash
```

```
# Check pod events
```

```
kubectl describe pod <POD-NAME>
```

```
# Check resource availability
```

```
kubectl top nodes
```

## 3. CNI Issues

```
bash
```

```
# Restart CNI pods
```

```
kubectl delete pods -n kube-system -l k8s-app=calico-node
```

```
# Check CNI Logs
```

```
kubectl logs -n kube-system -l k8s-app=calico-node
```

## 4. Reset Cluster (if needed)

```
bash
```

```
# Reset kubeadm
```

```
sudo kubeadm reset
```

```
sudo rm -rf /etc/cni/net.d
```

```
sudo rm -rf $HOME/.kube/config
```

---

## Security Considerations

### Production Hardening

- **Enable RBAC** (enabled by default in kubeadm)
- **Use Network Policies** for pod-to-pod communication

- **Regular security updates** for all components
- **Proper secret management** with tools like Sealed Secrets or External Secrets
- **Enable audit logging**
- **Use Pod Security Standards**

## Resource Limits

```
yaml

# Example resource limits
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-quota
spec:
  hard:
    requests.cpu: "1000m"
    requests.memory: 1Gi
    limits.cpu: "2000m"
    limits.memory: 2Gi
```

## Next Steps

1. **Set up monitoring** with Prometheus & Grafana
2. **Configure logging** with ELK stack or Loki
3. **Implement CI/CD** pipelines
4. **Set up backup strategies** for etcd and persistent volumes
5. **Plan for high availability** with multiple master nodes
6. **Implement security policies** and network segmentation

---

**Congratulations!** 🎉 You now have a fully functional Kubernetes cluster running with containerd. Your cluster is ready for deploying applications and learning advanced Kubernetes concepts.